


# CUDA Merge Sort

M.VENKATAKRISHNA

# MERGE SORT PARALLEL ALGORITHM (2-Stages)


EXPLAINED WITH 4 INITIAL SPLITS

# Stage-1


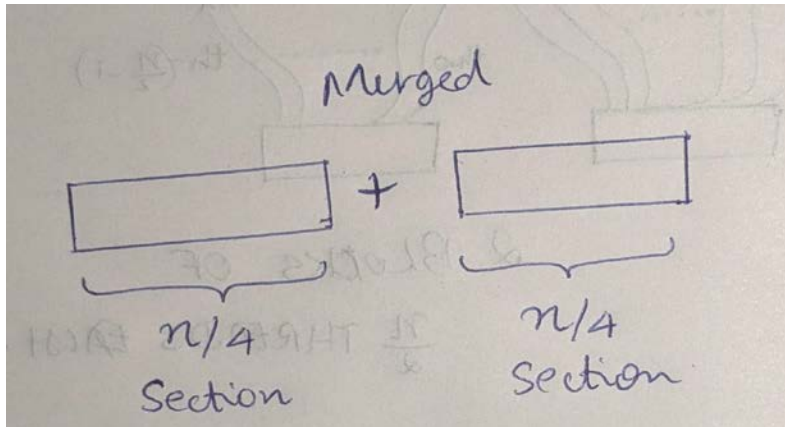


- The original array `arr[]` is of size  $n$
- The array `arr[]` is initially split into 4 sections
- A block of 4 threads is launched : **Each thread is responsible for sorting each of the 4 sections.** (BUBBLE-SORT IS EMPLOYED FOR THIS)

# STAGE-2 (Iteration-1)




- Now, each section of  $n/4$  elements has been sorted (STAGE-1)
- **2 blocks of  $n/2$  threads** are launched.
- Each block parallelly **merges a pair of 2 sections of size:  $n/4$  to produce a sorted SUPERSET of size:  $n/2$**
- Each Block of  $n/2$  threads is responsible for producing **1 SUPERSET of size:  $n/2$**







# STAGE-2 (Iteration-2)



- Now, each section of  $n/2$  elements has been sorted (STAGE-2 Iteration 1)
- **1 block of n threads** are launched.
- This block **merges a pair of 2 sections of size:  $n/2$  to produce a sorted SUPERSET of size:  $n$**
- This Block of n threads is responsible for producing **1 SUPERSET of size-n**




# MERGE SORT PARALLEL ALGORITHM


EXAMPLE RUN WITH 4 INITIAL SPLITS

Example `arr[]` = {5,1,15,14,10,13,3,2,20,17,21,22,18,16,25,24}

# STAGE-1: ARRAY IS DIVIDED INTO SECTIONS AND EACH SECTION IS SORTED





- **4 threads** are launched
- Each thread **individually sorts a section of 4 elements**





## STAGE 2: ITERATION-1

### EACH SORTED SECTION IS MERGED IN PAIRS TO PRODUCE SORTED SUPERSETS


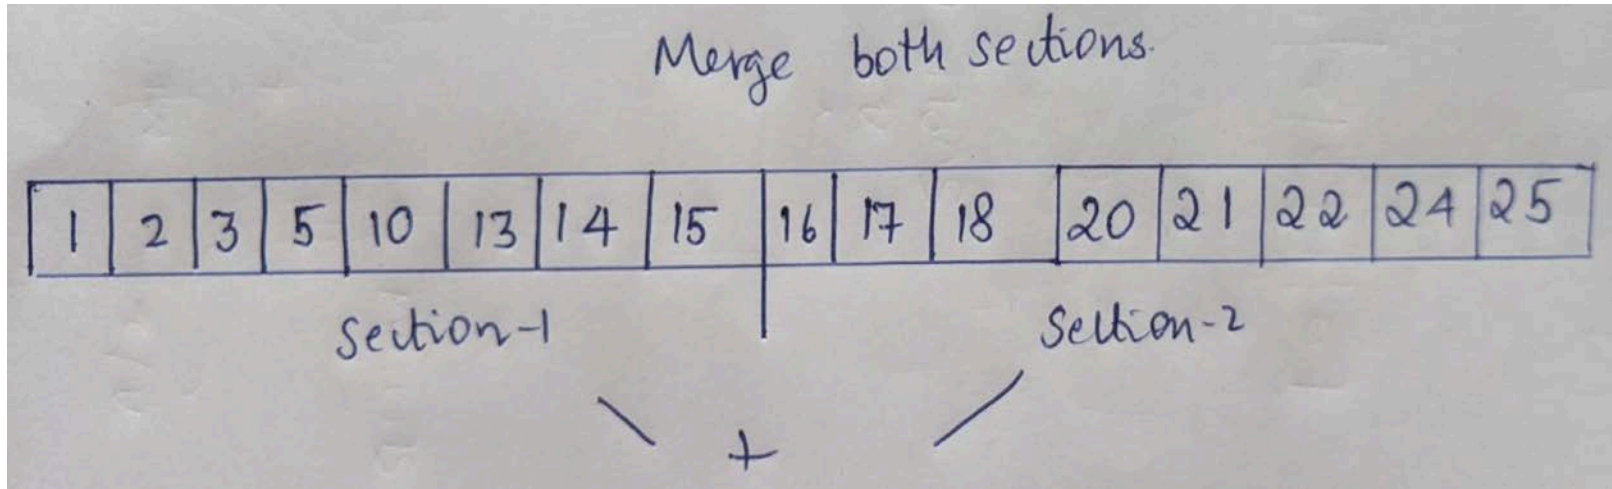


- 2 Sections on left are merged to form a sorted superset in left
- 2 sections in right are merged to form a sorted superset in right



## STAGE 2: ITERATION-2

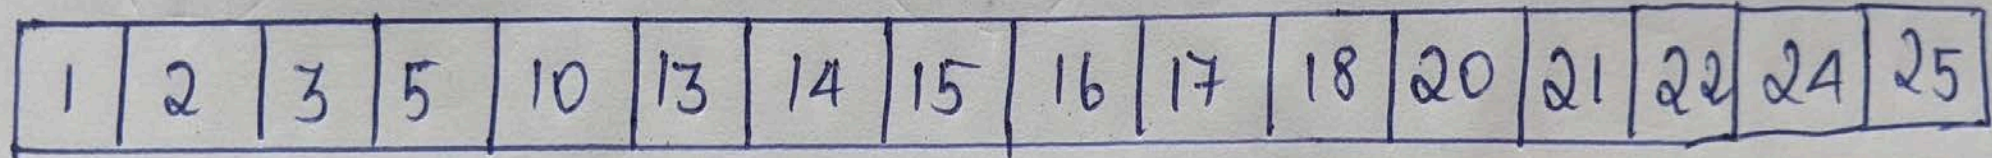
EACH SORTED SECTION IS MERGED IN PAIRS TO PRODUCE SORTED SUPERSETS



**2 Sorted Sections of size- $n/2$  each are merged to form a sorted superset of size  $n$**

# SORTED ARRAY IS FINALLY OBTAINED

**arr[] = {1,2,3,5,10,13,14,15,16,17,18,20,21,22,24,25}**

A hand-drawn diagram of an array. It consists of a horizontal row of 16 rectangular boxes, each containing a number. The numbers are 1, 2, 3, 5, 10, 13, 14, 15, 16, 17, 18, 20, 21, 22, 24, and 25, arranged in ascending order from left to right. The boxes are drawn with blue ink on a light-colored background.

1	2	3	5	10	13	14	15	16	17	18	20	21	22	24	25
---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----

# MERGE SORT CUDA IMPLEMENTATION

# KERNEL : section\_sort()

```
//Entire block for the entire array. Each thread takes care of bubble-sorting an individual section of size : *section_length
__global__ void section_sort(int* nums, int section_size)  //(, int n)
{
    //The thread with thread index = idx will take care of nums[] from : [( section_size * idx ) to ( section_size * (idx + 1) - 1 )]
    //For example: idx = 1 and section_size = 20, then, thread with idx = 1 will take care of nums[ 20: 39 ]

    int idx = threadIdx.x;

    //Bubble sort nums[] from index [ ( section_size * idx ) : ( section_size * (idx + 1) - 1 ) ]
    bubblesort(nums, section_size * idx, (section_size * (idx + 1) - 1));
}
```

## Takes as INPUT:

- `nums*` : The input array , which is to be divided into sections and sorted sectionally
- `section_size` : The size (number of elements) of each section

## Produces OUTPUT:

- `nums[]` Input array but with **each section individually sorted**



# section\_sort() : KERNEL-CALL AND WORKING

```
section_sort <<<1, div_num>>> (d_arr, section_size);
```

- *div\_num* = Number of sections into which arr[] is split into.
- A single block of *div\_num* number of threads are launched.
- Each thread will be responsible for sorting one specific section of the arr[].
- Thread number *idx* will be responsible for the section:  
[ *idx* \* *section\_size* : (*idx*+1) \* *section\_size* - 1 ]
- Bubble-Sort() is employed by each thread to sort the section its responsible for.

# Stage-2

Merge Sorted Sections

# Merge Kernel Call

```
//Number of Threads = Size of Superset  
//Number of blocks = Number of supersets  
merge <<< superset_num, n/superset_num, n/superset_num >>> (d_arr, section_size, d_out_temp);
```

- **Number of Blocks** = Number of Supersets
- **Number of threads** = Length of each superset
  
- **Each Block** : Is responsible for a superset (merging the 2 sections of its superset)
- **Each thread** : Is responsible for 1 element of the superset

# KERNEL: merge()

```
__global__ void merge(int* arr, int section_length, int* d_out_temp)
{
    //int section_length = *section_size;
    int superset_length = section_length * 2;    //Block will be 2 * (size of 1 section). Because 2 sections are merged
    int idx = threadIdx.x;
    int b_idx = blockIdx.x;

    //Length of arr1[] and arr2[] are section size
    int len1 = section_length;
    int len2 = section_length;

    //-----Select *arr1 and *arr2 and *d_out_curr-----

    int* arr1 = arr + (b_idx * superset_length);
    int* arr2 = arr1 + (section_length);
    int* d_out_curr = d_out_temp + (b_idx * superset_length);    //Determine d_out_curr[], the output array for current merge

    //Dynamically allocated shared memory array.
    // scat_ad[] from index [0 to n1-1] is for arr1[].
    //scat_ad[] from index [n1 to n2-1] is for arr2[]

    //Create a shared memory of size n1+n2 to accomodate the scatter-addresses corresponding to each element in arr1[] and arr2[]
    extern __shared__ int scat_ad[];
```

## Takes as INPUT:

- **arr\*** : Input array
- **section\_length** : length of each section, which are to be merged in pairs
- **d\_out\_temp\*** : temporary array used to store intermediate output

## Produces As Output:

- **arr[]** but with pairs of sections of **section\_length** merged to form bigger sections of size **section\_length\*2**



# Merge Kernel :Details

- **section** : Divisions in arr[] . They will be merged in pairs.
- **superset** : The result formed from merging to 2 sorted sections
- **superset\_length** : Length of each superset.
- **arr1\*** : Points to first section of superset.
- **arr2\*** : Points to second section of superset
- **len1 = len2** = Length of arr1[] = Length of arr2[]
- **scat\_ad[]** : **A shared memory array of length = superset-length.**
  - It will hold the scatter addresses for each element of sorted superset

# Merge : Scatter Address Calculation

```
//-----These threads are responsible for arr1[]-----  
if (idx <= len1 - 1)  
{  
    int idx1 = idx;    //Number of elements in arr1[] that are lesser than arr1[idx]. idx1 = index of current element in arr1[]  
  
    int target = arr1[idx1];    //Target is current element in arr1[]  
  
    //-----Find idx2-----Binary Search Part-----  
    int idx2 = bin_search(arr2, target, len2) + 1;    //Number of elements in arr2[] that are lesser than arr1[idx].....  
  
    //Calculate and store the scatter address in array  
    //scat_arr1[idx] = idx1 + idx2;    //If there are 2 elements before a number in output array, its index will be 2  
  
    scat_ad[idx] = idx1 + idx2;    //Scatter address corresponding to arr1[idx] = idx1 + idx2  
}
```

**Each thread is responsible for producing scatter address of one element in the superset**

- **If (thread index < len1):** Current thread is responsible for an element in arr1[]
- **Target = arr1[idx] (Current Number whose scatter address is to be calculated)**
- **Idx1** = Number of elements smaller than target in arr1[] = idx
- **Idx2** = Number of elements smaller than target in arr2[]
  - **To calculate idx2, we perform a binary\_search for target in arr2[]**
- **Scatter address of arr[idx] in sorted array will be (idx1 + idx2), i.e, the number of elements in superset smaller than target. Hence, store scatt\_ad[idx]= idx1+idx2**

```

//-----These threads are responsible for arr2[]-----
else if (idx >= len1)
{
    //Number of elements in arr2[] that are lesser than arr2[idx].
    //idx1 = index of current element in arr2[]
    //(idx-len1) because threads with index n1 to n2-1 are responsible for arr2[] index [0: n2-1]
    int idx1 = idx - len1;

    int target = arr2[idx1];    //Target is current element in arr1[]

    //-----Find idx2-----Binary Search Part-----
    int idx2 = bin_search(arr1, target, len1) + 1;    //Number of elements in arr1[] that are lesser than arr2[idx]. +1 bcos we want appro

    //Calculate and store the scatter address in array
    //scat_arr1[idx] = idx1 + idx2;    //If there are 2 elements before a number in output array, its index will be 2
    scat_ad[idx] = idx1 + idx2;    //Scatter address corresponding to arr2[idx - len1] = idx1 + idx2
}

__syncthreads();    //Barrier to ensure that all threads have finished writing scat_ad[]-----Not necessary

```

**If idx >= len1:** Current thread is responsible for an element in arr2[]

Calculating scatter address is similar to previous case

**\_\_syncthread()** : Barrier to ensure all threads have finished calculating scatter addresses.

# Place the elements in Superset in appropriate sorted positions

```
//-----Store the output in respective position in d_out_temp[] using scatter address so that they are in sorted order-----  
/* ... */  
d_out_curr[scat_ad[idx]] = arr1[idx];  
  
__syncthreads();  
//-----Copy sorted elements back to array-----  
  
arr1[idx] = d_out_curr[idx];  
  
//printf( "%d ", arr1[idx] );  
}
```

- 1) Store the elements in **current superset** from **arr[]** to **d\_out[]** in **appropriate sorted position** using **scatter address** from **scat\_ad[]** array (d\_out[] is temporary array)
- 2) Barrier : Ensure all threads have finished copying to d\_out[]
- 3) Copy back to arr[] from d\_out[]



# merge\_sort()

**CPU CODE**

# merge\_sort() Function

```
// ...
void merge_sort()
{
    GpuTimer timer;

    // ...

    int h_arr[64] = { 100,99,98,97,96,95,94,93,92,91,90,89,88,87,86,85,84,83,82,81,80,79,78,77,76,75,74,73,72,71,70,69,68,67,66,65,64,63,62,61

    int n = sizeof(h_arr) / sizeof(int);    //n = Total size of host array

    int div_num = 8;    //How many parts the array is initially split.
    int section_size = n/div_num;    // section_size = Size of each section after splitting arr[] into div_num parts (Stored in Hos

    //-----Create input and output arrays in GPU-----
    int* d_arr, * d_out_temp;    // *d_out2;
    cudaMalloc((void**)&d_arr, n * sizeof(int));
    cudaMemcpy((void*)d_arr, (void*)h_arr, n * sizeof(int), cudaMemcpyHostToDevice);    //d_arr[] is input array in device

    cudaMalloc((void**)&d_out_temp, n * sizeof(int));    //d_out_temp[] is temporarily used to store sorted block elements
```

- h\_arr[] : Input array to be sorted (Stored in HOST)
- d\_arr[] : Input array (Stored in Device)
- d\_out\_temp[] : temporary array (Device)

contd.

- **div\_num: Number of sections** into which the arr[] is initially SPLIT
- **section\_size:** Length of each section = Total array length / number of sections =  $n / \text{div\_num}$

## STAGE – 1: SECTION-SORT KERNEL CALL

```
//-----Stage-1: KERNEL CALL: Bubble Sort Each Section of section_size elements-----  
section_sort <<<1, div_num>>> (d_arr, section_size);    //Call div_num threads: Each thread bubble-sorts a sub-section of n/div_num el
```

- section-sort kernel call will bubble-sort all sections in arr[]. (The array is divided into div\_num number of sections each of section\_size)
- KERNEL calls 1 block of div\_num threads, 1 thread for each section in arr[].
- Each thread will sort its section



# STAGE-2: ITERATIVELY MERGE SECTIONS UNTIL ENTIRE ARRAY IS SORTED

```
//Initially, section_size = n / div_num
int superset_num = div_num / 2; //Is the total number of supersets in the array, each of which are merged by a separate block. Initially,
while (superset_num >= 1)
{
    /* ... */
    // ...

    //Number of Threads = Size of Superset
    //Number of blocks = Number of supersets
    merge <<< superset_num, n/superset_num, n/superset_num >>> (d_arr, section_size, d_out_temp);

    //UPDATE : superset_num ( halved ) and section_size (doubled)
    superset_num = superset_num/2;
    section_size = section_size*2;
}
```

- Iteratively call merge kernel
- In each iteration, merge sections into supersets.
- After each iteration, **section size is doubled and number of superset is halved** (Because after each iteration, each pair of section is merged, so that current superset is a section for the next iteration )
- **In Final iteration, only a single pair of sections is to be merged and then, sorted array will be obtained**

END